

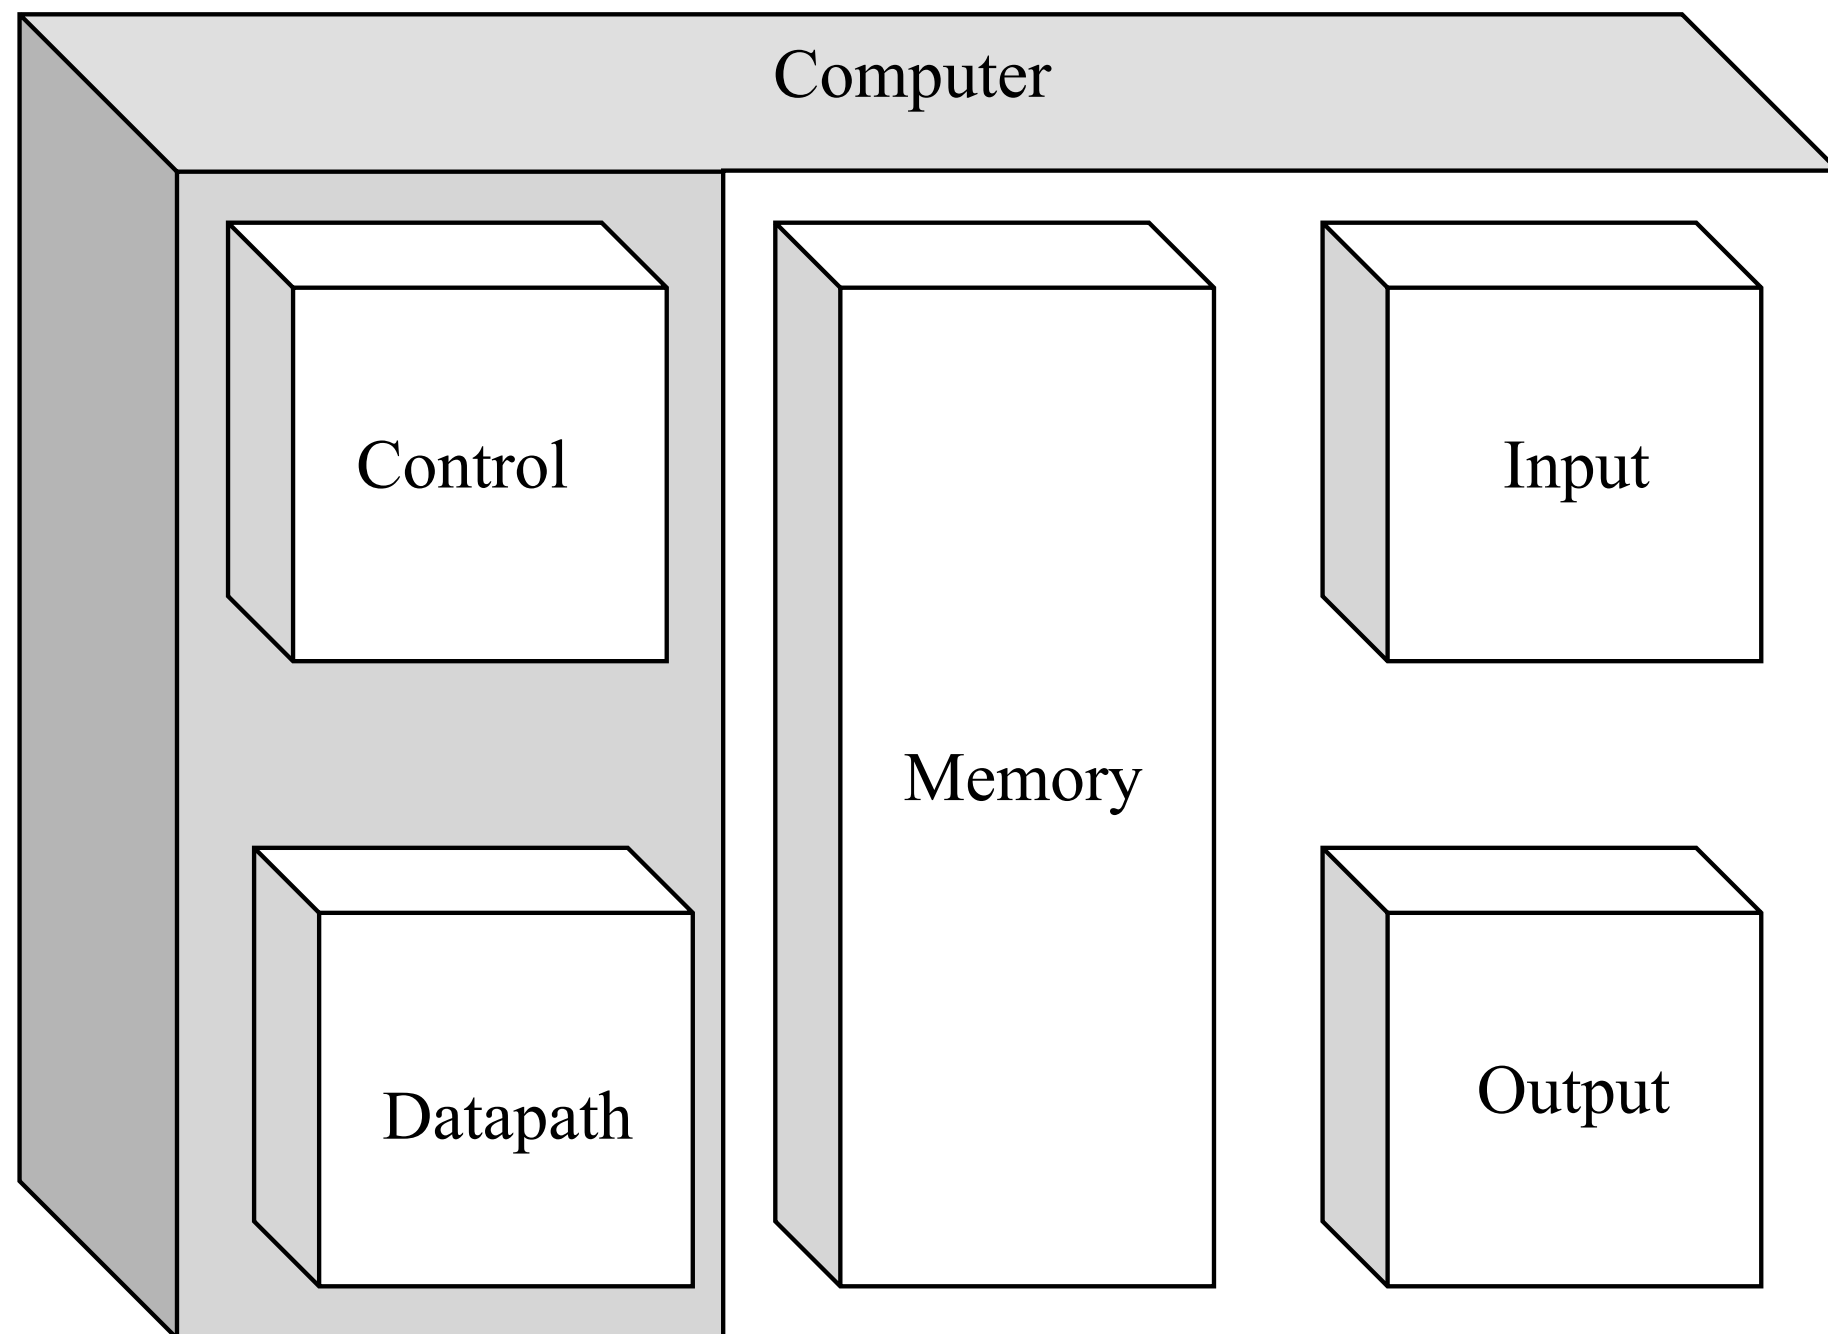
# Memory Subsystem Design

---

Jason Mars

# The Memory Subsystem

---



# Memory Locality

---

# Memory Locality

---

- Memory hierarchies take advantage of memory locality.

# Memory Locality

---

- Memory hierarchies take advantage of memory locality.
- Memory locality is the principle that future memory accesses are near past accesses.

# Memory Locality

---

- Memory hierarchies take advantage of memory locality.
- Memory locality is the principle that future memory accesses are near past accesses.
- Memories take advantage of two types of locality
  - -- near in time => we will often access the same data again very soon
  - -- near in space/distance => our next access is often very close to our last access (or recent accesses).

# Memory Locality

---

- Memory hierarchies take advantage of memory locality.
- Memory locality is the principle that future memory accesses are near past accesses.
- Memories take advantage of two types of locality
  - **Temporal Locality** -- near in time => we will often access the same data again very soon
  - -- near in space/distance => our next access is often very close to our last access (or recent accesses).

# Memory Locality

---

- Memory hierarchies take advantage of memory locality.
- Memory locality is the principle that future memory accesses are near past accesses.
- Memories take advantage of two types of locality
  - **Temporal Locality** -- near in time => we will often access the same data again very soon
  - **Spacial Locality** -- near in space/distance => our next access is often very close to our last access (or recent accesses).



# Memory Locality

---

- Memory hierarchies take advantage of memory locality.
- Memory locality is the principle that future memory accesses are near past accesses.
- Memories take advantage of two types of locality
  - **Temporal Locality** -- near in time => we will often access the same data again very soon
  - **Spacial Locality** -- near in space/distance => our next access is often very close to our last access (or recent accesses).

(this sequence of addresses exhibits both temporal and spatial locality)

1,2,3,1,2,3,8,8,47,9,10,8,8...

# Memory Locality

---

- Memory hierarchies take advantage of memory locality.
- Memory locality is the principle that future memory accesses are near past accesses.
- Memories take advantage of two types of locality
  - **Temporal Locality** -- near in time => we will often access the same data again very soon
  - **Spacial Locality** -- near in space/distance => our next access is often very close to our last access (or recent accesses).

(this sequence of addresses exhibits both temporal and spatial locality)

1,2,3,1,2,3,8,8,47,9,10,8,8...

# Memory Locality

---

- Memory hierarchies take advantage of memory locality.
- Memory locality is the principle that future memory accesses are near past accesses.
- Memories take advantage of two types of locality
  - **Temporal Locality** -- near in time => we will often access the same data again very soon
  - **Spacial Locality** -- near in space/distance => our next access is often very close to our last access (or recent accesses).

(this sequence of addresses exhibits both temporal and spatial locality)

1,2,3,1,2,3,8,8,47,9,10,8,8...

# Memory Locality

---

- Memory hierarchies take advantage of memory locality.
- Memory locality is the principle that future memory accesses are near past accesses.
- Memories take advantage of two types of locality
  - **Temporal Locality** -- near in time => we will often access the same data again very soon
  - **Spacial Locality** -- near in space/distance => our next access is often very close to our last access (or recent accesses).

(this sequence of addresses exhibits both temporal and spatial locality)

1,2,3,1,2,3,8,8,47,9,10,8,8...

# Locality and Cacheing

---

- Memory hierarchies exploit locality by cacheing (keeping close to the processor) data likely to be used again.
- This is done because we can build large, slow memories and small, fast memories, but we can't build large, fast memories.
- If it works, we get the illusion of SRAM access time with disk capacity

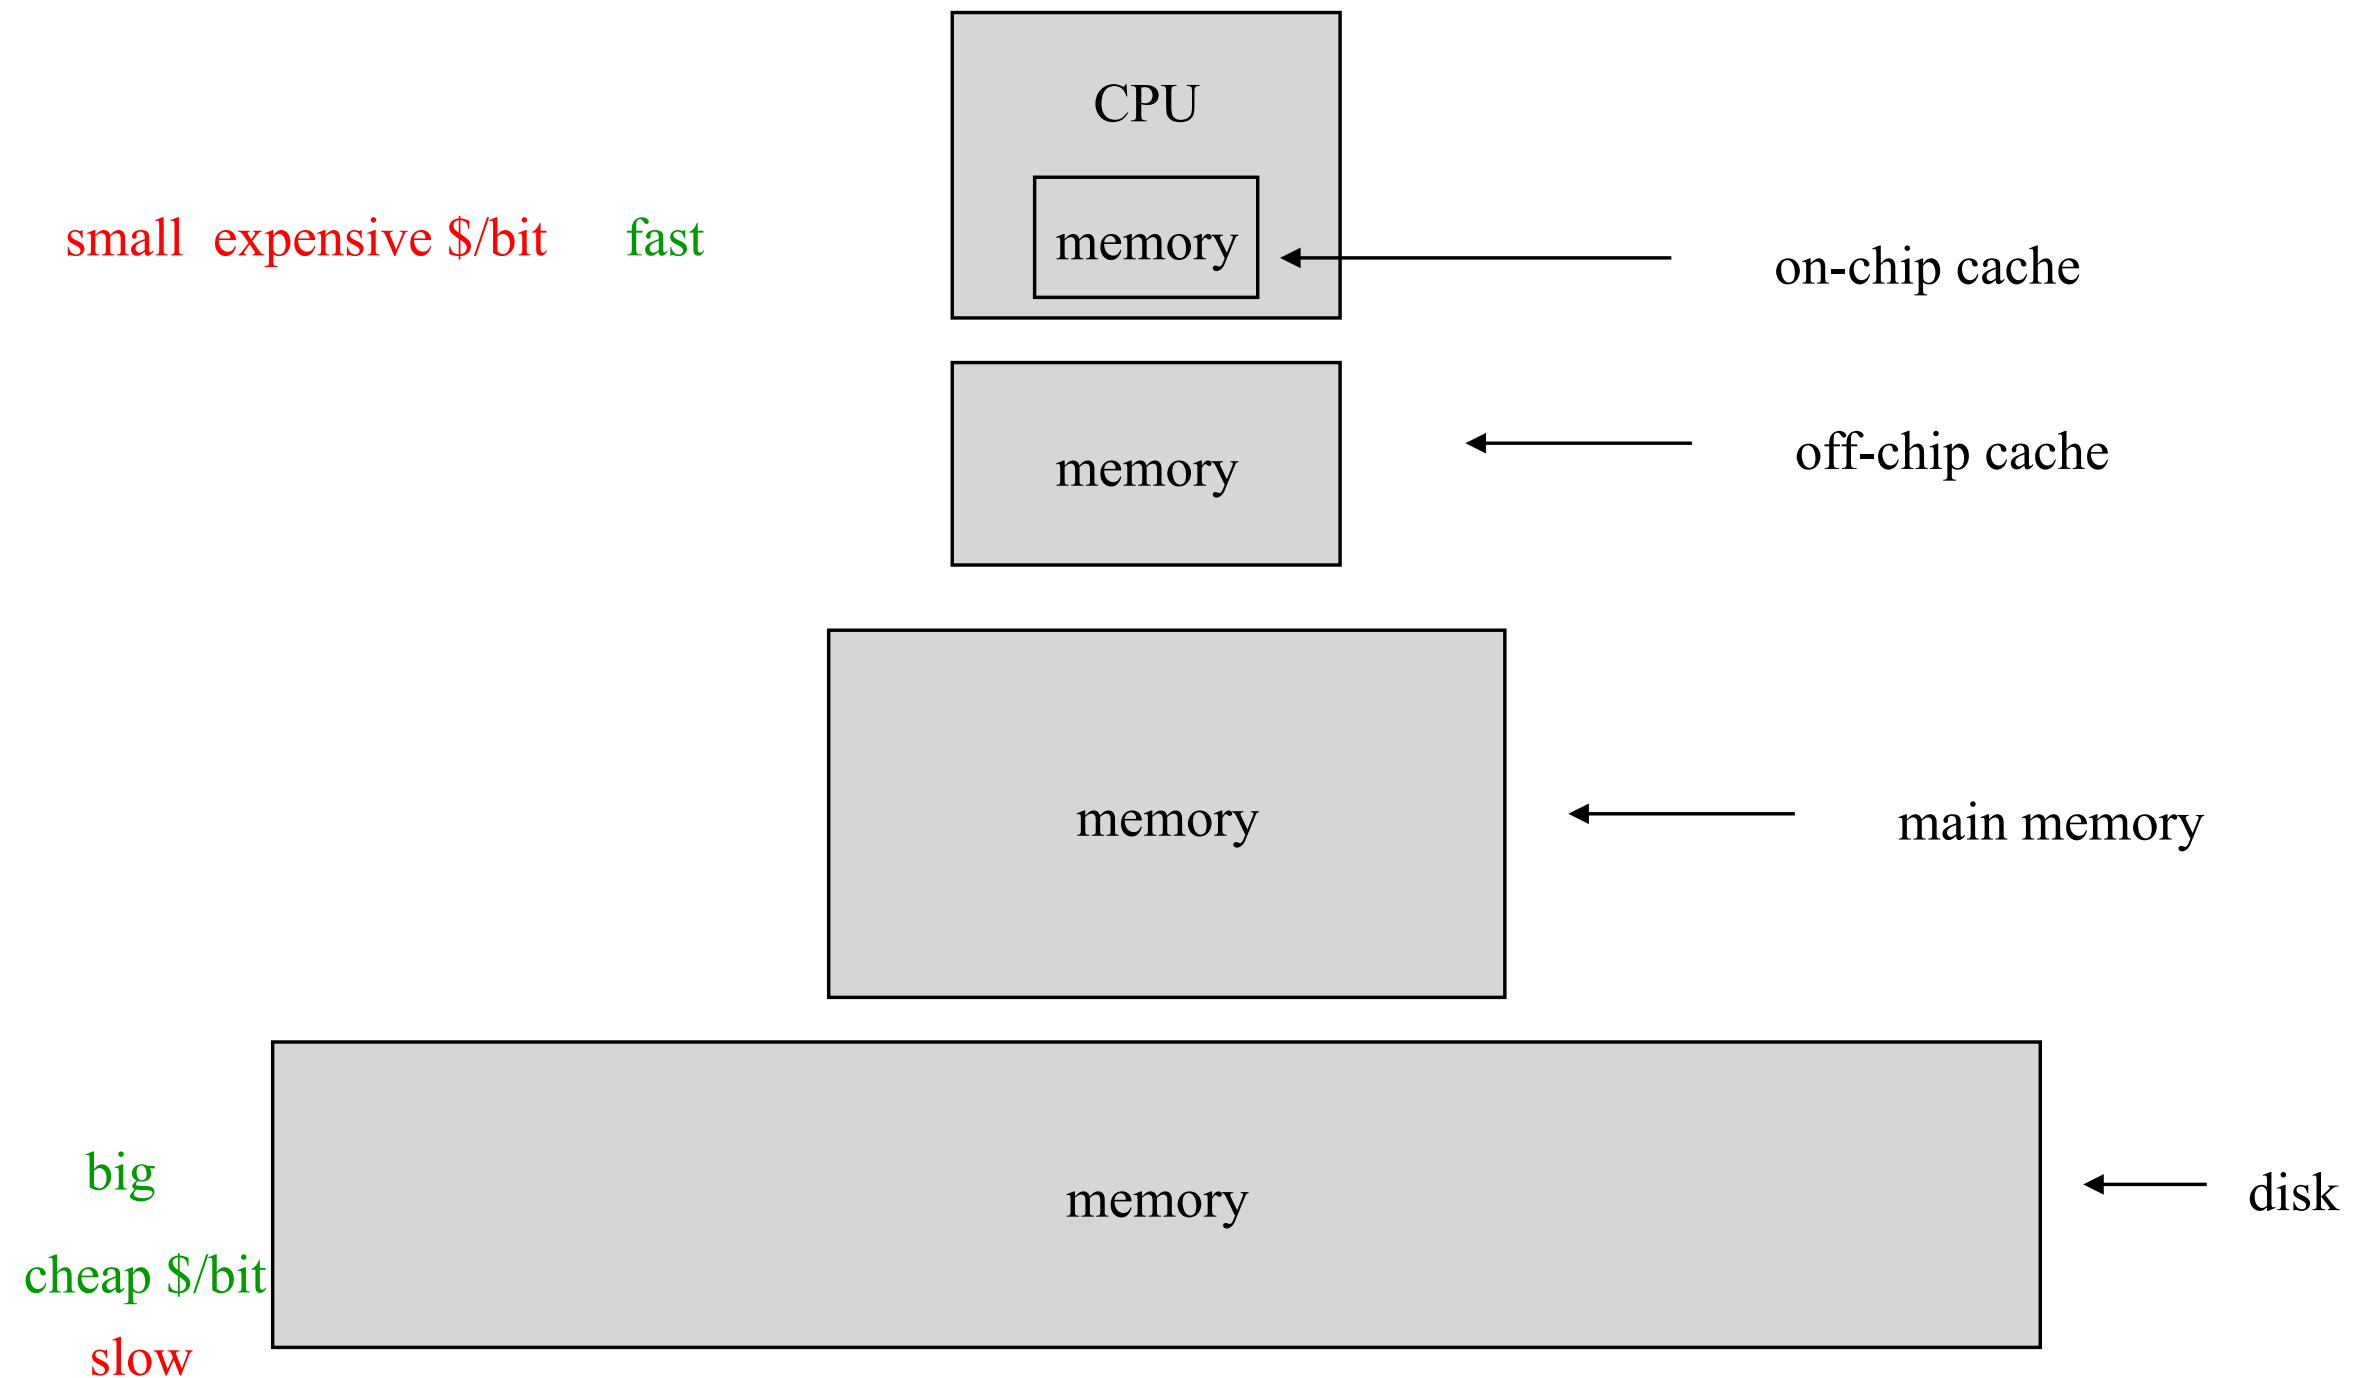
**SRAM access times are ~1ns at cost of \$2000 to \$5000 per Gbyte.**

**DRAM access times are ~70ns at cost of \$20 to \$75 per Gbyte.**

**Disk access times are 5 to 20 million ns at cost of \$.20 to \$2 per Gbyte.**


# Typical Memory Hierarchy

---




# Cache Fundamentals

---



# Cache Fundamentals

---




*cache hit* -- an access where the data is found in the cache.



# Cache Fundamentals

---




*cache hit* -- an access where the data is found in the cache.

*cache miss* -- an access which isn't

# Cache Fundamentals

---




*cache hit* -- an access where the data is found in the cache.

*cache miss* -- an access which isn't

*hit time* -- time to access the cache

# Cache Fundamentals

---



*cache hit* -- an access where the data is found in the cache.


*cache miss* -- an access which isn't

*hit time* -- time to access the cache

*hit ratio* -- percentage of time the data is found in the cache

# Cache Fundamentals

---



*cache hit* -- an access where the data is found in the cache.

*cache miss* -- an access which isn't


*hit time* -- time to access the cache

*hit ratio* -- percentage of time the data is found in the cache

*miss penalty* -- time to move data from further level to closer, then to cpu

# Cache Fundamentals

---



***cache hit*** -- an access where the data is found in the cache.

***cache miss*** -- an access which isn't

***hit time*** -- time to access the cache


***hit ratio*** -- percentage of time the data is found in the cache

***miss penalty*** -- time to move data from further level to closer, then to cpu

***miss ratio*** --  $(1 - \text{hit ratio})$

# More Fundamentals


---



# More Fundamentals

---


- **cache block size** or **cache line size** -- the amount of data that gets transferred on a cache miss.



# More Fundamentals

---

- **cache block size** or **cache line size** -- the amount of data that gets transferred on a cache miss.
- **instruction cache** -- cache that only holds instructions.






# More Fundamentals

---


- **cache block size** or **cache line size** -- the amount of data that gets transferred on a cache miss.
- **instruction cache** -- cache that only holds instructions.
- **data cache** -- cache that only caches data.



# More Fundamentals

---


- **cache block size** or **cache line size** -- the amount of data that gets transferred on a cache miss.
- **instruction cache** -- cache that only holds instructions.
- **data cache** -- cache that only caches data.
- **unified cache** -- cache that holds both.



# Caching Issues

---

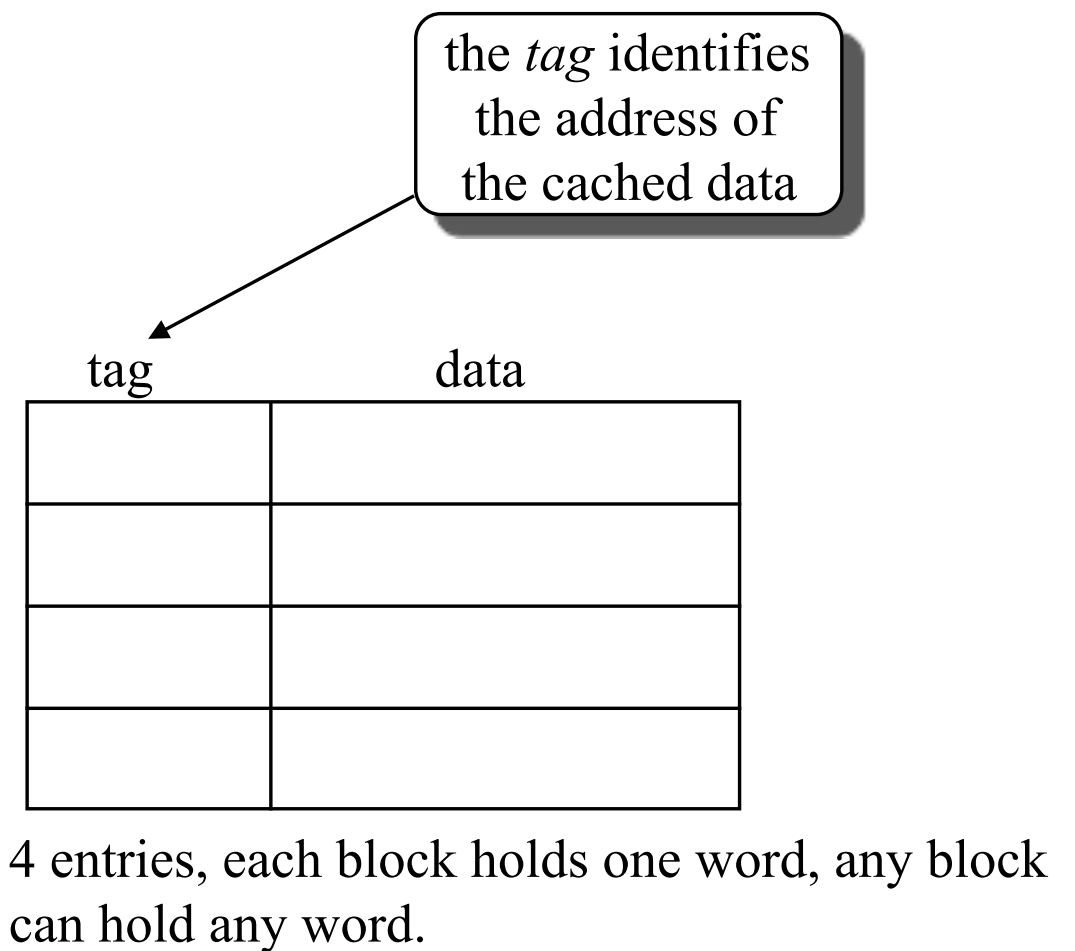
- On a memory access -
  - How do I know if this is a hit or miss?
- On a cache miss -
  - where to put the new data?
  - what data to throw out?
  - how to remember what data this is?



# A Simple Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100



- A cache that can put a line of data anywhere is called \_\_\_\_\_
- The most popular replacement strategy is LRU ( \_\_\_\_\_ ).

# A Simple Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

the *tag* identifies  
the address of  
the cached data

tag	data

4 entries, each block holds one word, any block can hold any word.

- A cache that can put a line of data anywhere is called **Fully Associative**
- The most popular replacement strategy is LRU ( ).

# A Simple Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

the *tag* identifies  
the address of  
the cached data

tag	data

4 entries, each block holds one word, any block can hold any word.

- A cache that can put a line of data anywhere is called **Fully Associative**
- The most popular replacement strategy is LRU ( ).

# A Simple Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

the *tag* identifies  
the address of  
the cached data

tag	data
000001	

4 entries, each block holds one word, any block can hold any word.

- A cache that can put a line of data anywhere is called Fully Associative
- The most popular replacement strategy is LRU ( ).

# A Simple Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

the *tag* identifies  
the address of  
the cached data

tag	data
000001	Mem(000001)

4 entries, each block holds one word, any block can hold any word.

- A cache that can put a line of data anywhere is called Fully Associative
- The most popular replacement strategy is LRU ( ).



# A Simple Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

the *tag* identifies  
the address of  
the cached data

tag	data
000001	Mem(000001)
000010	Mem(000010)

4 entries, each block holds one word, any block can hold any word.

- A cache that can put a line of data anywhere is called **Fully Associative**
- The most popular replacement strategy is LRU ( ).

# A Simple Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

the *tag* identifies  
the address of  
the cached data

tag	data
000001	Mem(000001)
000010	Mem(000010)
000011	Mem(000011)

4 entries, each block holds one word, any block can hold any word.

- A cache that can put a line of data anywhere is called **Fully Associative**
- The most popular replacement strategy is LRU ( ).

# A Simple Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

the *tag* identifies  
the address of  
the cached data

tag	data
000001	Mem(000001)
000010	Mem(000010)
000011	Mem(000011)

4 entries, each block holds one word, any block can hold any word.

- A cache that can put a line of data anywhere is called **Fully Associative**
- The most popular replacement strategy is LRU ( ).

# A Simple Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

the *tag* identifies  
the address of  
the cached data

tag	data
000001	Mem(000001)
000010	Mem(000010)
000011	Mem(000011)

4 entries, each block holds one word, any block can hold any word.

- A cache that can put a line of data anywhere is called **Fully Associative**
- The most popular replacement strategy is LRU ( ).

# A Simple Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

the *tag* identifies  
the address of  
the cached data

tag	data
000001	Mem(000001)
000010	Mem(000010)
000011	Mem(000011)

4 entries, each block holds one word, any block can hold any word.

- A cache that can put a line of data anywhere is called **Fully Associative**
- The most popular replacement strategy is LRU ( ).

# A Simple Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

the *tag* identifies  
the address of  
the cached data

tag	data
000001	Mem(000001)
000010	Mem(000010)
000011	Mem(000011)
000101	Mem(000101)

4 entries, each block holds one word, any block can hold any word.

- A cache that can put a line of data anywhere is called **Fully Associative**
- The most popular replacement strategy is LRU ( ).

# A Simple Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

the *tag* identifies  
the address of  
the cached data

tag	data
000001	Mem(000001)
000010	Mem(000010)
000011	Mem(000011)
000101	Mem(000101)

4 entries, each block holds one word, any block can hold any word.

- A cache that can put a line of data anywhere is called **Fully Associative**
- The most popular replacement strategy is LRU ( ).

# A Simple Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

the *tag* identifies  
the address of  
the cached data

tag	data
000001	Mem(000001)
000010	Mem(000010)
000011	Mem(000011)
000101	Mem(000101)

4 entries, each block holds one word, any block can hold any word.

- A cache that can put a line of data anywhere is called **Fully Associative**
- The most popular replacement strategy is LRU ( ).



# A Simple Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

the *tag* identifies  
the address of  
the cached data

tag	data
000001	Mem(000001)
000010	Mem(000010)
000011	Mem(000011)
000101	Mem(000101)

4 entries, each block holds one word, any block can hold any word.

- A cache that can put a line of data anywhere is called **Fully Associative**
- The most popular replacement strategy is LRU (**Least Recently Used**).

# A Simple Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

the *tag* identifies  
the address of  
the cached data

tag	data
000001	Mem(000001)
000010	Mem(000010)
000011	Mem(000011)
000101	Mem(000101)

4 entries, each block holds one word, any block can hold any word.

- A cache that can put a line of data anywhere is called **Fully Associative**
- The most popular replacement strategy is LRU (**Least Recently Used**).

# A Simple Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

the *tag* identifies  
the address of  
the cached data

tag	data
000001	Mem(000001)
000010	Mem(000010)
000011	Mem(000011)
000101	Mem(000101)

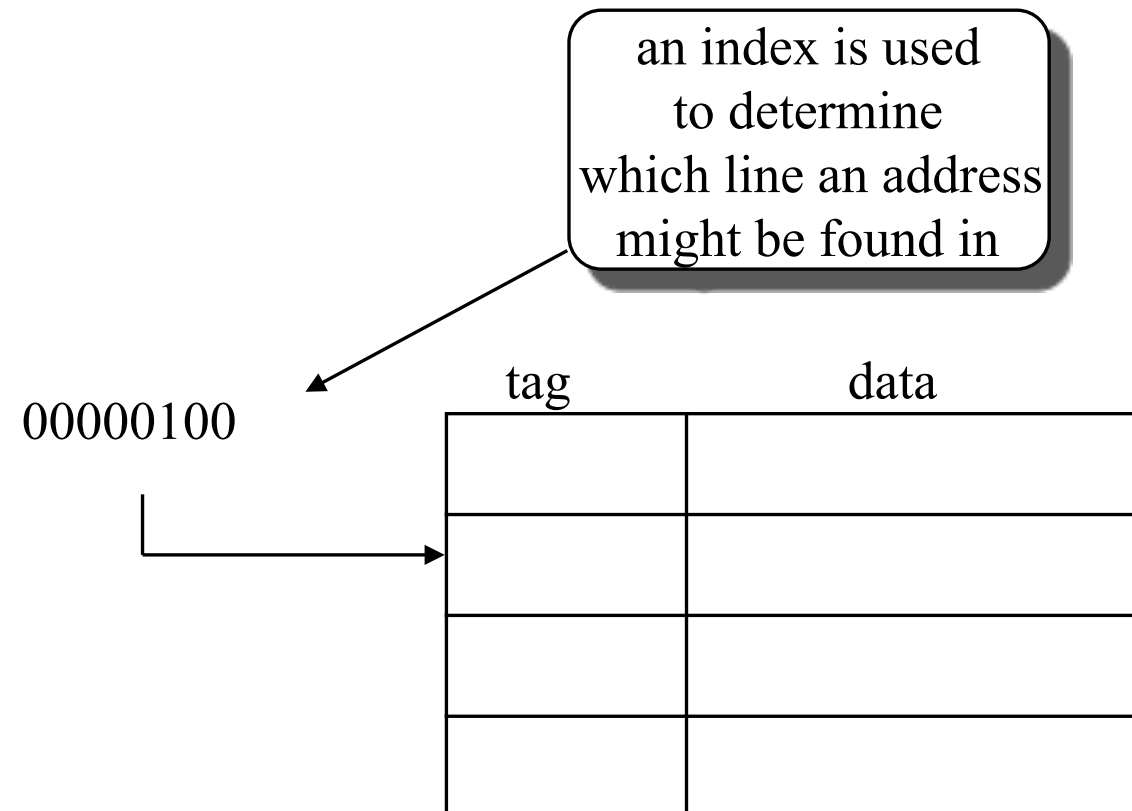
4 entries, each block holds one word, any block can hold any word.

- A cache that can put a line of data anywhere is called **Fully Associative**
- The most popular replacement strategy is LRU (**Least Recently Used**).

# An Even Simpler Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100




4 entries, each block holds one word, each word in memory maps to exactly one cache location.

- A cache that can put a line of data in exactly one place is called \_\_\_\_\_.
- Advantages/disadvantages vs. fully-associative?

# An Even Simpler Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100




4 entries, each block holds one word, each word in memory maps to exactly one cache location.

- A cache that can put a line of data in exactly one place is called **Direct Mapped**.
- Advantages/disadvantages vs. fully-associative?

# An Even Simpler Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100



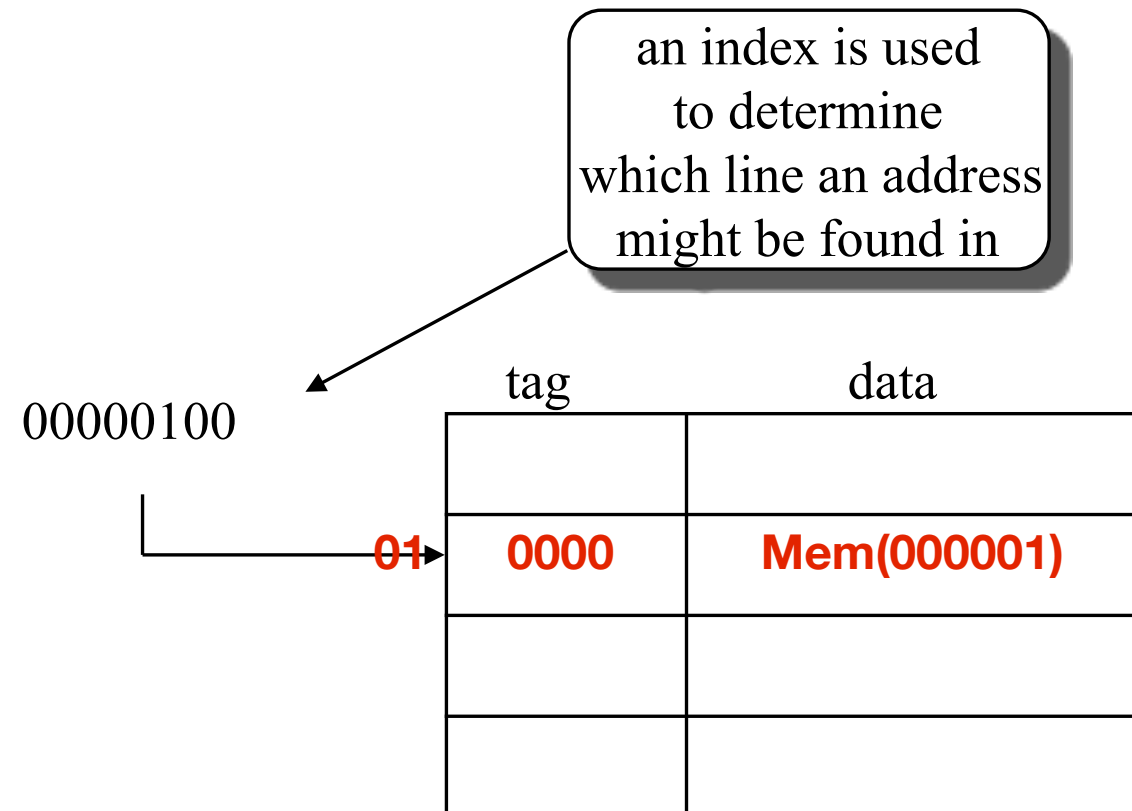
4 entries, each block holds one word, each word in memory maps to exactly one cache location.

- A cache that can put a line of data in exactly one place is called **Direct Mapped**.
- Advantages/disadvantages vs. fully-associative?

# An Even Simpler Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100



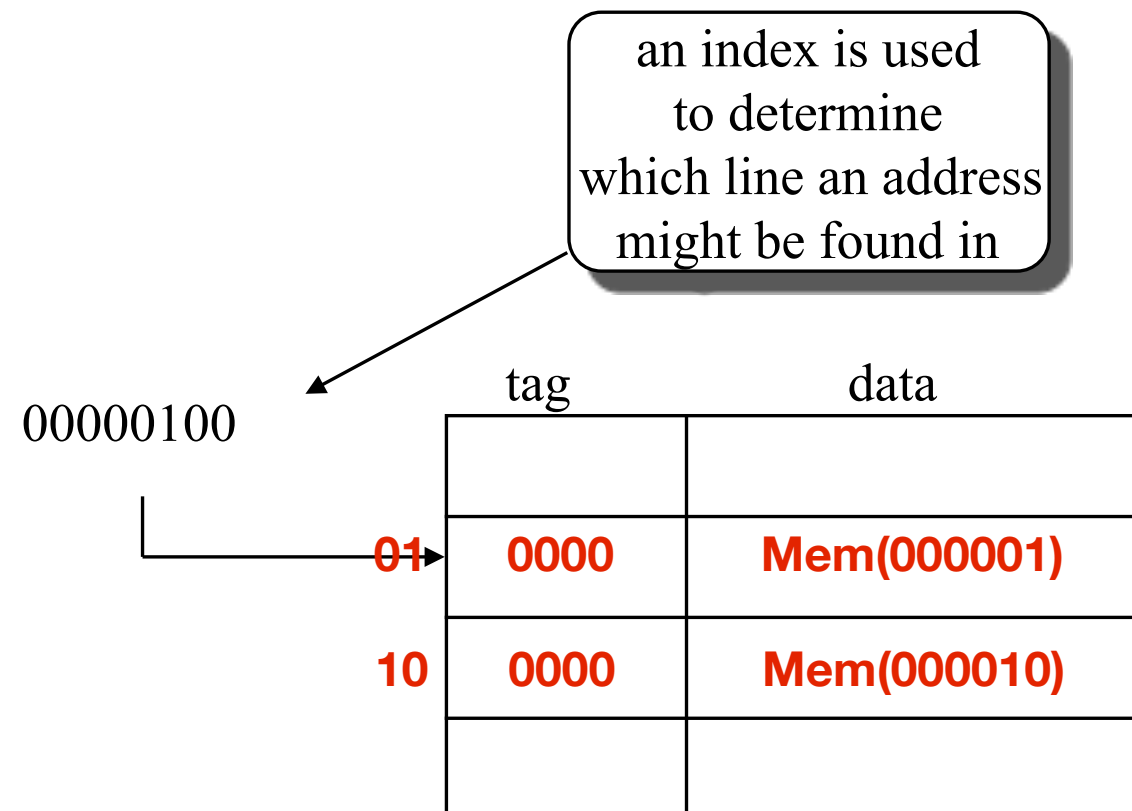
4 entries, each block holds one word, each word in memory maps to exactly one cache location.

- A cache that can put a line of data in exactly one place is called **Direct Mapped**.
- Advantages/disadvantages vs. fully-associative?

# An Even Simpler Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100



4 entries, each block holds one word, each word in memory maps to exactly one cache location.

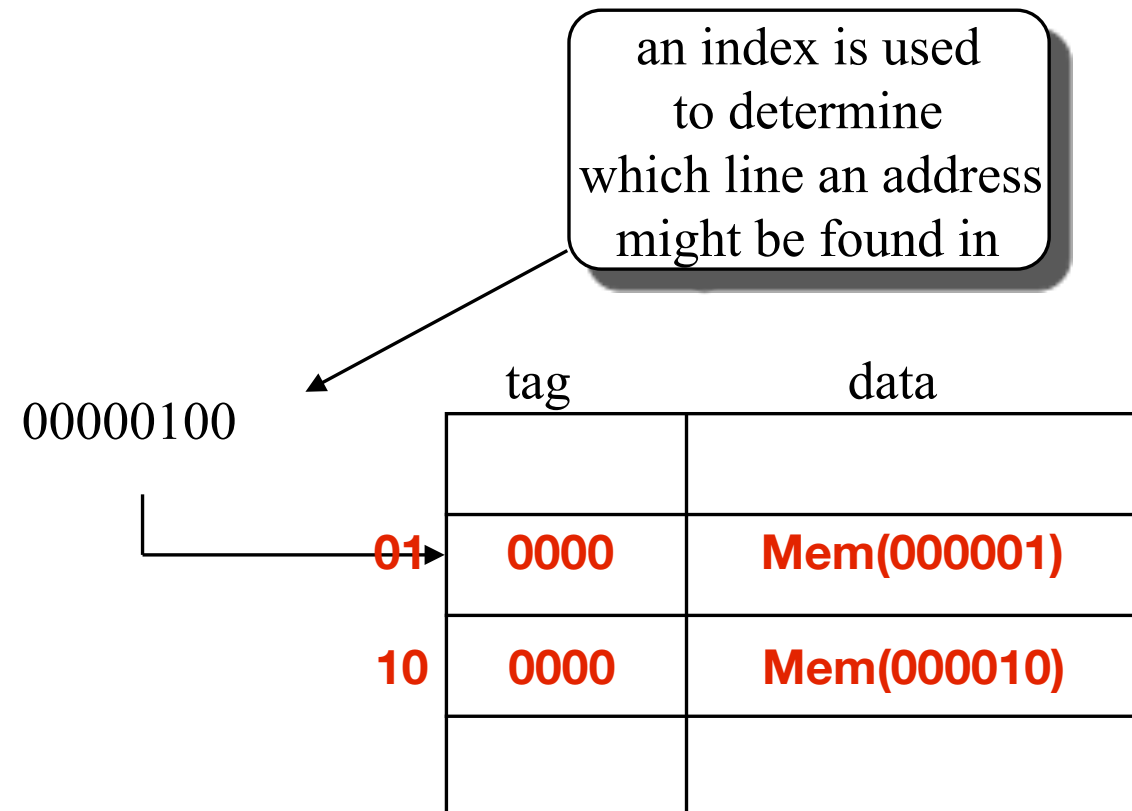
- A cache that can put a line of data in exactly one place is called **Direct Mapped**.
- Advantages/disadvantages vs. fully-associative?



# An Even Simpler Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100



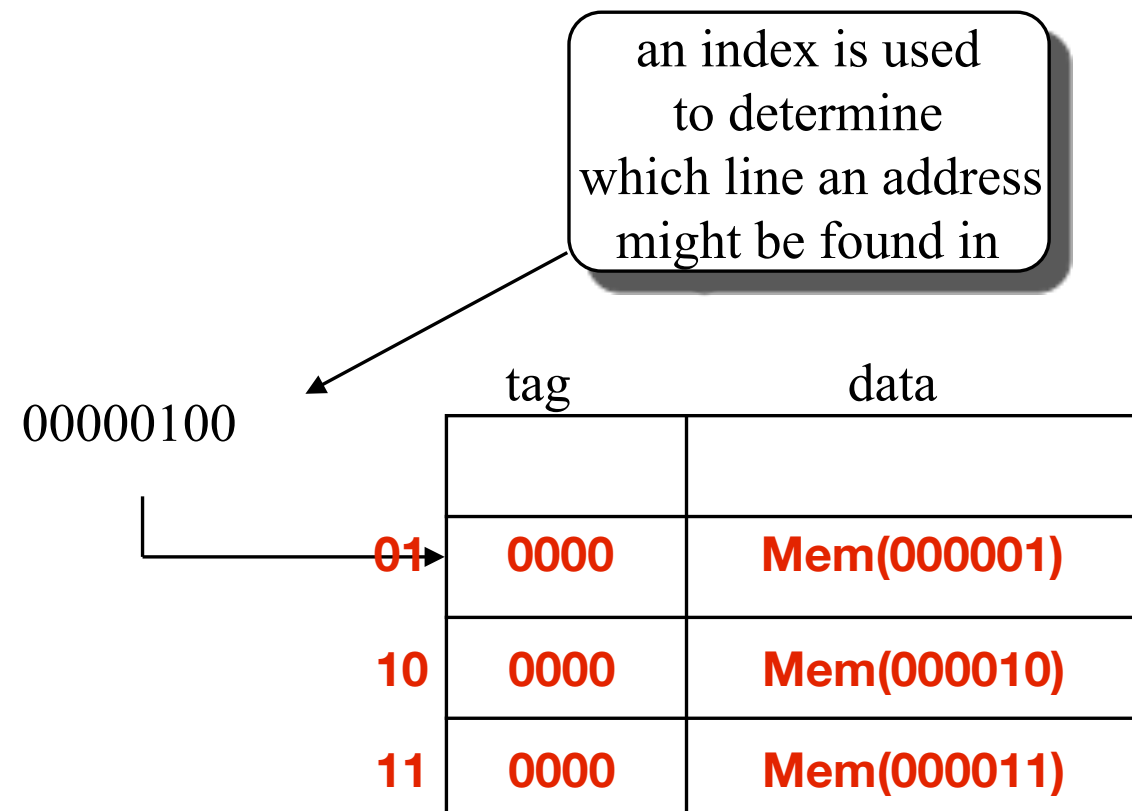
4 entries, each block holds one word, each word in memory maps to exactly one cache location.

- A cache that can put a line of data in exactly one place is called **Direct Mapped**.
- Advantages/disadvantages vs. fully-associative?

# An Even Simpler Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100




4 entries, each block holds one word, each word in memory maps to exactly one cache location.

- A cache that can put a line of data in exactly one place is called **Direct Mapped**.
- Advantages/disadvantages vs. fully-associative?

# An Even Simpler Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100




4 entries, each block holds one word, each word in memory maps to exactly one cache location.

- A cache that can put a line of data in exactly one place is called **Direct Mapped**.
- Advantages/disadvantages vs. fully-associative?

# An Even Simpler Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100



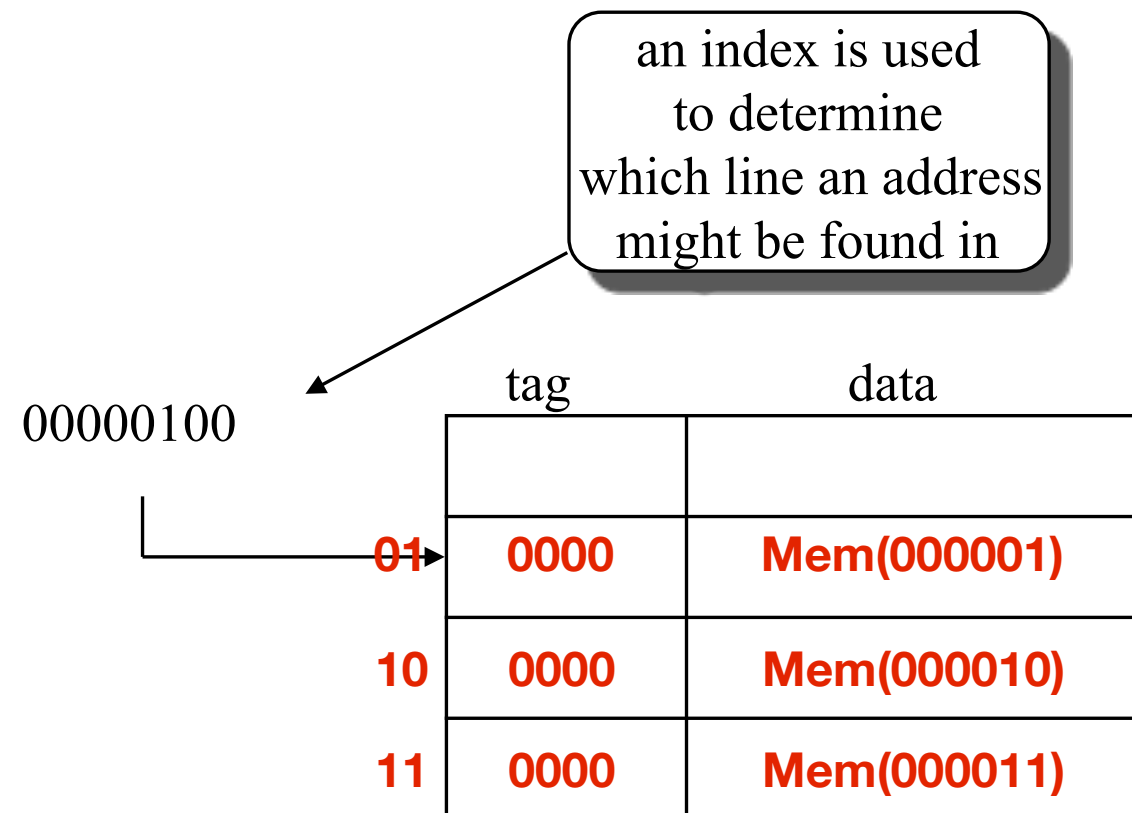
4 entries, each block holds one word, each word in memory maps to exactly one cache location.

- A cache that can put a line of data in exactly one place is called **Direct Mapped**.
- Advantages/disadvantages vs. fully-associative?

# An Even Simpler Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100



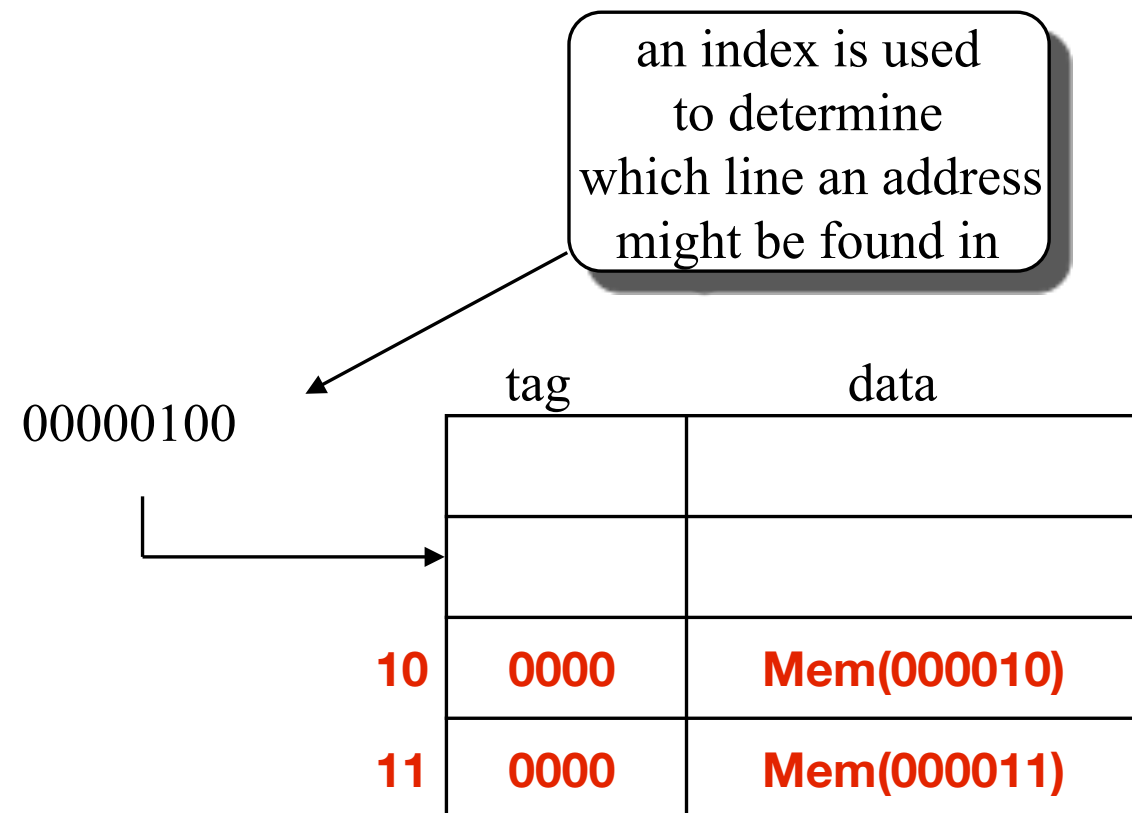
4 entries, each block holds one word, each word in memory maps to exactly one cache location.

- A cache that can put a line of data in exactly one place is called **Direct Mapped**.
- Advantages/disadvantages vs. fully-associative?

# An Even Simpler Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100



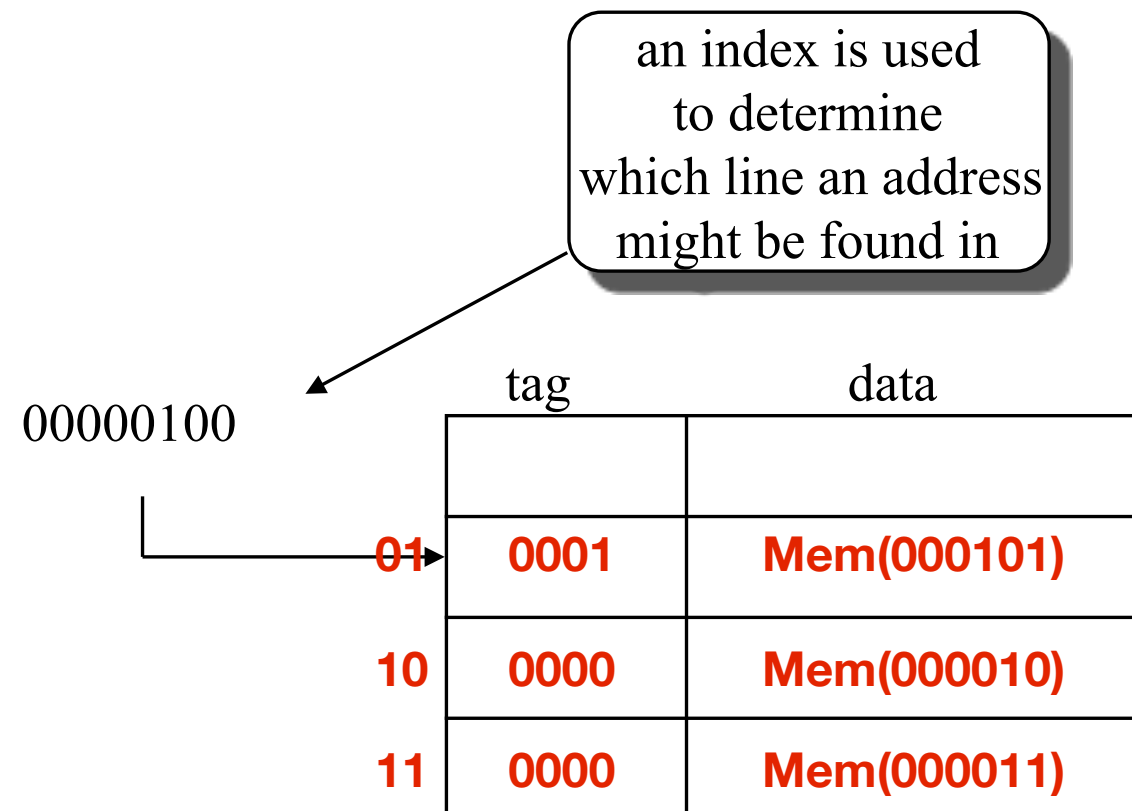
4 entries, each block holds one word, each word in memory maps to exactly one cache location.

- A cache that can put a line of data in exactly one place is called **Direct Mapped**.
- Advantages/disadvantages vs. fully-associative?

# An Even Simpler Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100



4 entries, each block holds one word, each word in memory maps to exactly one cache location.


- A cache that can put a line of data in exactly one place is called **Direct Mapped**.
- Advantages/disadvantages vs. fully-associative?

# A Set Associative Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

00000100



4 entries, each block holds one word, each word in memory maps to one of a set of  $n$  cache lines

- A cache that can put a line of data in exactly  $n$  places is called \_\_\_\_\_.
- The cache lines/blocks that share the same index are a cache \_\_\_\_\_.




# A Set Associative Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

00000100



4 entries, each block holds one word, each word in memory maps to one of a set of  $n$  cache lines


- A cache that can put a line of data in exactly  $n$  places is called  **$n$ -way set-associative**.
- The cache lines/blocks that share the same index are a cache \_\_\_\_\_.

# A Set Associative Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

00000100



4 entries, each block holds one word, each word in memory maps to one of a set of  $n$  cache lines


- A cache that can put a line of data in exactly  $n$  places is called **n-way set-associative**.
- The cache lines/blocks that share the same index are a cache **set**.

# A Set Associative Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

00000100



4 entries, each block holds one word, each word in memory maps to one of a set of  $n$  cache lines

- A cache that can put a line of data in exactly  $n$  places is called  **$n$ -way set-associative**.
- The cache lines/blocks that share the same index are a cache **set**.

# A Set Associative Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

00000100

	tag	data	tag	data
0				
1	00000	Mem(000001)		

4 entries, each block holds one word, each word in memory maps to one of a set of  $n$  cache lines

- A cache that can put a line of data in exactly  $n$  places is called  **$n$ -way set-associative**.
- The cache lines/blocks that share the same index are a cache **set**.

# A Set Associative Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

00000100

	tag	data	tag	data
0				
1	00000	Mem(000001)		

4 entries, each block holds one word, each word in memory maps to one of a set of  $n$  cache lines

- A cache that can put a line of data in exactly  $n$  places is called **n-way set-associative**.
- The cache lines/blocks that share the same index are a cache **set**.

# A Set Associative Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

00000100

	tag	data	tag	data
0	00001	Mem(000010)		
1	00000	Mem(000001)		

4 entries, each block holds one word, each word in memory maps to one of a set of  $n$  cache lines

- A cache that can put a line of data in exactly  $n$  places is called  **$n$ -way set-associative**.
- The cache lines/blocks that share the same index are a cache **set**.

# A Set Associative Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

00000100

	tag	data	tag	data
0	00001	Mem(000010)		
1	00000	Mem(000001)		

4 entries, each block holds one word, each word in memory maps to one of a set of  $n$  cache lines

- A cache that can put a line of data in exactly  $n$  places is called  **$n$ -way set-associative**.
- The cache lines/blocks that share the same index are a cache **set**.

# A Set Associative Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

00000100

	tag	data	tag	data
0	00001	Mem(000010)		
1	00000	Mem(000001)	00001	Mem(000011)

4 entries, each block holds one word, each word in memory maps to one of a set of  $n$  cache lines

- A cache that can put a line of data in exactly  $n$  places is called  **$n$ -way set-associative**.
- The cache lines/blocks that share the same index are a cache **set**.



# A Set Associative Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

00000100

	tag	data	tag	data
0	00001	Mem(000010)		
1	00000	Mem(000001)	00001	Mem(000011)

4 entries, each block holds one word, each word in memory maps to one of a set of  $n$  cache lines

- A cache that can put a line of data in exactly  $n$  places is called  **$n$ -way set-associative**.
- The cache lines/blocks that share the same index are a cache **set**.

# A Set Associative Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

00000100

	tag	data	tag	data
0	00001	Mem(000010)		
1	00000	Mem(000001)	00001	Mem(000011)

4 entries, each block holds one word, each word in memory maps to one of a set of  $n$  cache lines

- A cache that can put a line of data in exactly  $n$  places is called  **$n$ -way set-associative**.
- The cache lines/blocks that share the same index are a cache **set**.

# A Set Associative Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

00000100

	tag	data	tag	data
0	00001	Mem(000010)		
1	00000	Mem(000001)	00001	Mem(000011)

4 entries, each block holds one word, each word in memory maps to one of a set of  $n$  cache lines

- A cache that can put a line of data in exactly  $n$  places is called  **$n$ -way set-associative**.
- The cache lines/blocks that share the same index are a cache **set**.

# A Set Associative Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

00000100

	tag	data	tag	data
0	00001	Mem(000010)		
1			00001	Mem(000011)

4 entries, each block holds one word, each word in memory maps to one of a set of  $n$  cache lines

- A cache that can put a line of data in exactly  $n$  places is called  **$n$ -way set-associative**.
- The cache lines/blocks that share the same index are a cache **set**.

# A Set Associative Cache

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

00000100

	tag	data	tag	data
0	00001	Mem(000010)		
1	00010	Mem(000101)	00001	Mem(000011)

4 entries, each block holds one word, each word in memory maps to one of a set of  $n$  cache lines

- A cache that can put a line of data in exactly  $n$  places is called  **$n$ -way set-associative**.
- The cache lines/blocks that share the same index are a cache **set**.

# Longer/Larger Cache Blocks

---

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

00000100



tag	data	

4 entries, each block holds two words, each word in memory maps to exactly one cache location (this cache is twice the total size of the prior caches).

- Large cache blocks take advantage of \_\_\_\_\_.
- Too large of a block size can waste cache space.
- Longer cache blocks require less tag space

# Longer/Larger Cache Blocks

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

00000100



tag	data	

4 entries, each block holds two words, each word in memory maps to exactly one cache location (this cache is twice the total size of the prior caches).

- Large cache blocks take advantage of Spacial Locality.
- Too large of a block size can waste cache space.
- Longer cache blocks require less tag space

# Longer/Larger Cache Blocks

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

00000100

00  
01  
10  
11

tag

data

tag	data

4 entries, each block holds two words, each word in memory maps to exactly one cache location (this cache is twice the total size of the prior caches).

- Large cache blocks take advantage of Spacial Locality.
- Too large of a block size can waste cache space.
- Longer cache blocks require less tag space



# Longer/Larger Cache Blocks

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

00000100

00  
01  
10  
11

tag

0

data


4 entries, each block holds two words, each word in memory maps to exactly one cache location (this cache is twice the total size of the prior caches).

- Large cache blocks take advantage of Spacial Locality.
- Too large of a block size can waste cache space.
- Longer cache blocks require less tag space

# Longer/Larger Cache Blocks

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

00000100

00  
01  
10  
11

tag

0

data 1


4 entries, each block holds two words, each word in memory maps to exactly one cache location (this cache is twice the total size of the prior caches).

- Large cache blocks take advantage of Spacial Locality.
- Too large of a block size can waste cache space.
- Longer cache blocks require less tag space

# Longer/Larger Cache Blocks

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

00000100

00

01

10

11

tag

0

data 1


4 entries, each block holds two words, each word in memory maps to exactly one cache location (this cache is twice the total size of the prior caches).

- Large cache blocks take advantage of Spacial Locality.
- Too large of a block size can waste cache space.
- Longer cache blocks require less tag space

# Longer/Larger Cache Blocks

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

00000100

00  
01  
10  
11

tag

0

data 1

tag	data
000	

4 entries, each block holds two words, each word in memory maps to exactly one cache location (this cache is twice the total size of the prior caches).

- Large cache blocks take advantage of Spacial Locality.
- Too large of a block size can waste cache space.
- Longer cache blocks require less tag space

# Longer/Larger Cache Blocks

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

00000100

00  
01  
10  
11

tag

0

data

1

tag	data
000	Mem(000000) Mem(000001)

4 entries, each block holds two words, each word in memory maps to exactly one cache location (this cache is twice the total size of the prior caches).

- Large cache blocks take advantage of Spacial Locality.
- Too large of a block size can waste cache space.
- Longer cache blocks require less tag space

# Longer/Larger Cache Blocks

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

00000100

00  
01  
10  
11

tag

0

data 1

tag	data
000	Mem(000000) Mem(000001)

4 entries, each block holds two words, each word in memory maps to exactly one cache location (this cache is twice the total size of the prior caches).

- Large cache blocks take advantage of Spacial Locality.
- Too large of a block size can waste cache space.
- Longer cache blocks require less tag space

# Longer/Larger Cache Blocks

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

00000100

00  
01  
10  
11

tag

0

data

1

tag	data
000	Mem(000000) Mem(000001)
000	

4 entries, each block holds two words, each word in memory maps to exactly one cache location (this cache is twice the total size of the prior caches).

- Large cache blocks take advantage of Spacial Locality.
- Too large of a block size can waste cache space.
- Longer cache blocks require less tag space

# Longer/Larger Cache Blocks

address string:

4	00000100
8	00001000
12	00001100
4	00000100
8	00001000
20	00010100
4	00000100
8	00001000
20	00010100
24	00011000
12	00001100
8	00001000
4	00000100

00000100

00  
01  
10  
11

tag

0

data

1


000	Mem(000000)	Mem(000001)
000	Mem(000010)	Mem(000011)

4 entries, each block holds two words, each word in memory maps to exactly one cache location (this cache is twice the total size of the prior caches).

- Large cache blocks take advantage of Spacial Locality.
- Too large of a block size can waste cache space.
- Longer cache blocks require less tag space



# Block Size and Miss Rate



# Cache Parameters

---

Cache size = Number of sets \* block size \* associativity

-128 blocks, 32-byte block size, direct mapped, size =

-128 KB cache, 64-byte blocks, 512 sets, associativity = ?

# Cache Parameters

---

Cache size = Number of sets \* block size \* associativity

-128 blocks, 32-byte block size, direct mapped, size =

$$128 \times 32 = 4096 \text{ bytes} = 4\text{mb}$$

-128 KB cache, 64-byte blocks, 512 sets, associativity = ?

# Cache Parameters

---

Cache size = Number of sets \* block size \* associativity

-128 blocks, 32-byte block size, direct mapped, size =

$$128 \times 32 = 4096 \text{ bytes} = 4\text{mb}$$

-128 KB cache, 64-byte blocks, 512 sets, associativity = ?

$$131072 \text{ bytes} / 512 = 256 \text{ bytes/set}$$

# Cache Parameters

---

Cache size = Number of sets \* block size \* associativity

-128 blocks, 32-byte block size, direct mapped, size =

$$128 \times 32 = 4096 \text{ bytes} = 4\text{mb}$$

-128 KB cache, 64-byte blocks, 512 sets, associativity = ?


$$131072 \text{ bytes} / 512 = 256 \text{ bytes/set}$$

$$256 \text{ bytes} / 64 \text{ byte} = 4 \text{ blocks/set} = 4\text{-way}$$

# A Cache Access


---

- 1. Use index and tag to access cache and determine hit/miss.
- 2. If hit, return requested data.
- 3. If miss, select a cache block to be replaced, and access memory or next lower cache (possibly stalling the processor).
  - load entire missed cache line into cache
  - return requested data to CPU (or higher cache)
- 4. If next lower memory is a cache, goto step 1 for that cache.



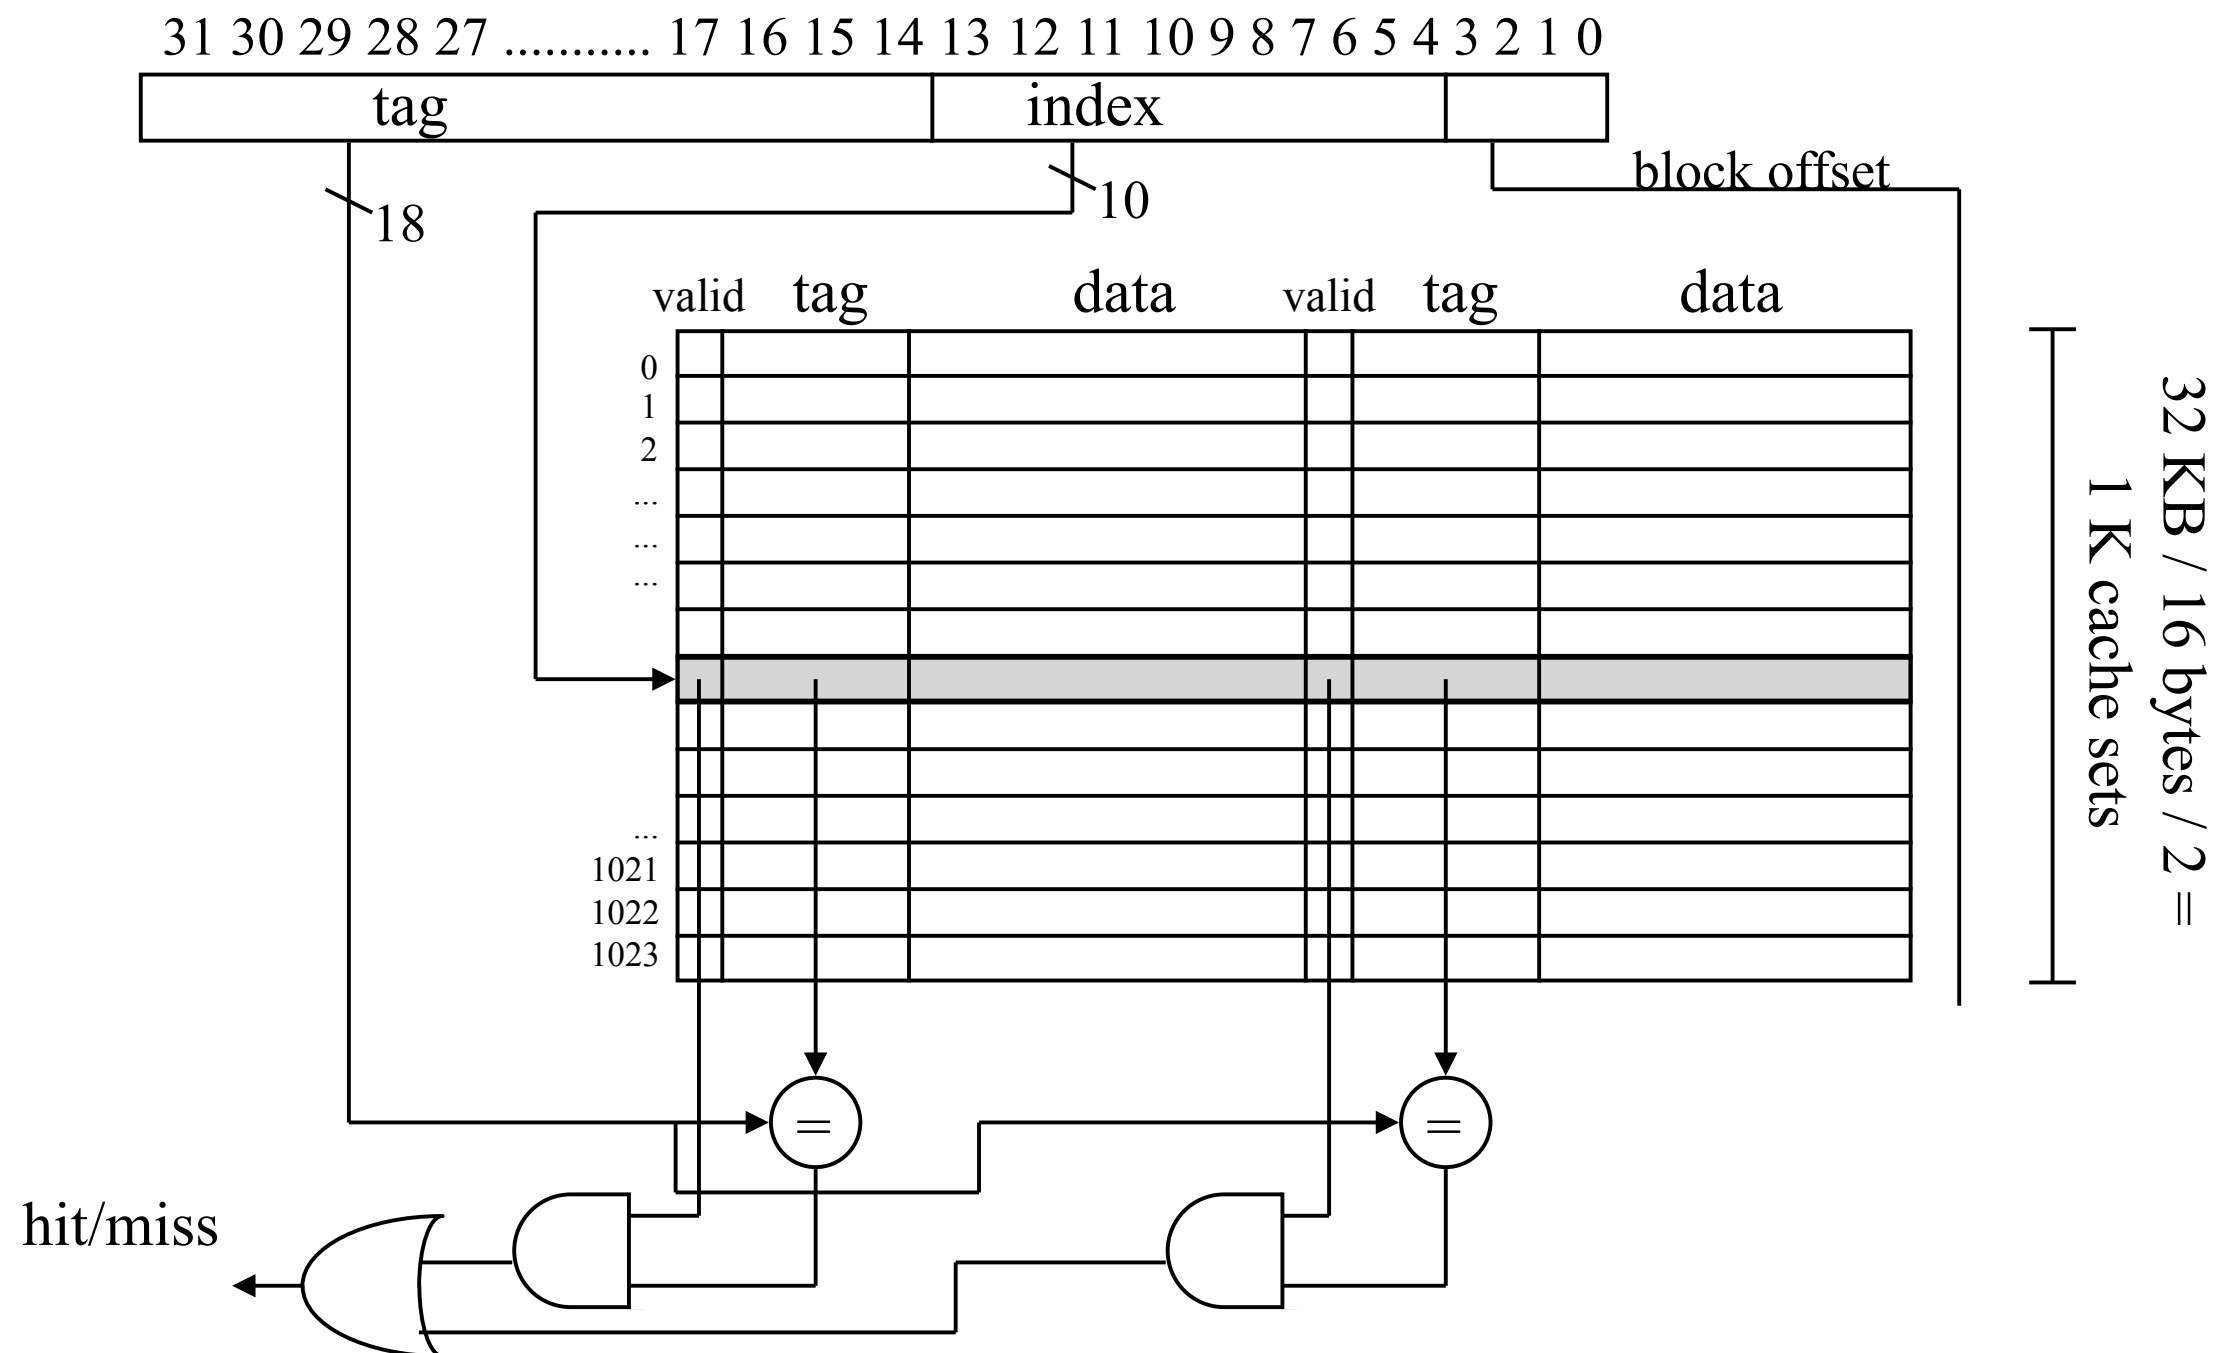
# Accessing a Sample Cache

- 16 KB cache, direct-mapped, 64-byte cache block size



# Accessing a Sample Cache

- 32 KB cache, 2-way set-associative, 16-byte block size





# Associative Caches

---

- Higher hit rates, but...
- longer access time (longer to determine hit/miss, more muxing of outputs)
- more space (longer tags compared to DM)
  - 2-way extra 1 bit
  - 4-way extra 2 bits

# Handling Stores

---

- Keep memory and cache identical?
- => all writes go to both cache and main memory
- => writes go only to cache. Modified cache lines are written back to memory when the line is replaced.
- Make room in cache for store miss?
- => on a store miss, bring written line into the cache
- => on a store miss, ignore cache

# Handling Stores

---

- Keep memory and cache identical?
- **Write-through** => all writes go to both cache and main memory
- => writes go only to cache. Modified cache lines are written back to memory when the line is replaced.
- Make room in cache for store miss?
- => on a store miss, bring written line into the cache
- => on a store miss, ignore cache

# Handling Stores

---

- Keep memory and cache identical?
- **Write-through** => all writes go to both cache and main memory
- **Write-back** => writes go only to cache. Modified cache lines are written back to memory when the line is replaced.
- Make room in cache for store miss?
- => on a store miss, bring written line into the cache
- => on a store miss, ignore cache

# Handling Stores

---

- Keep memory and cache identical?
- **Write-through** => all writes go to both cache and main memory
- **Write-back** => writes go only to cache. Modified cache lines are written back to memory when the line is replaced.
- Make room in cache for store miss?
- **Write-allocate** => on a store miss, bring written line into the cache
- => on a store miss, ignore cache

# Handling Stores

---

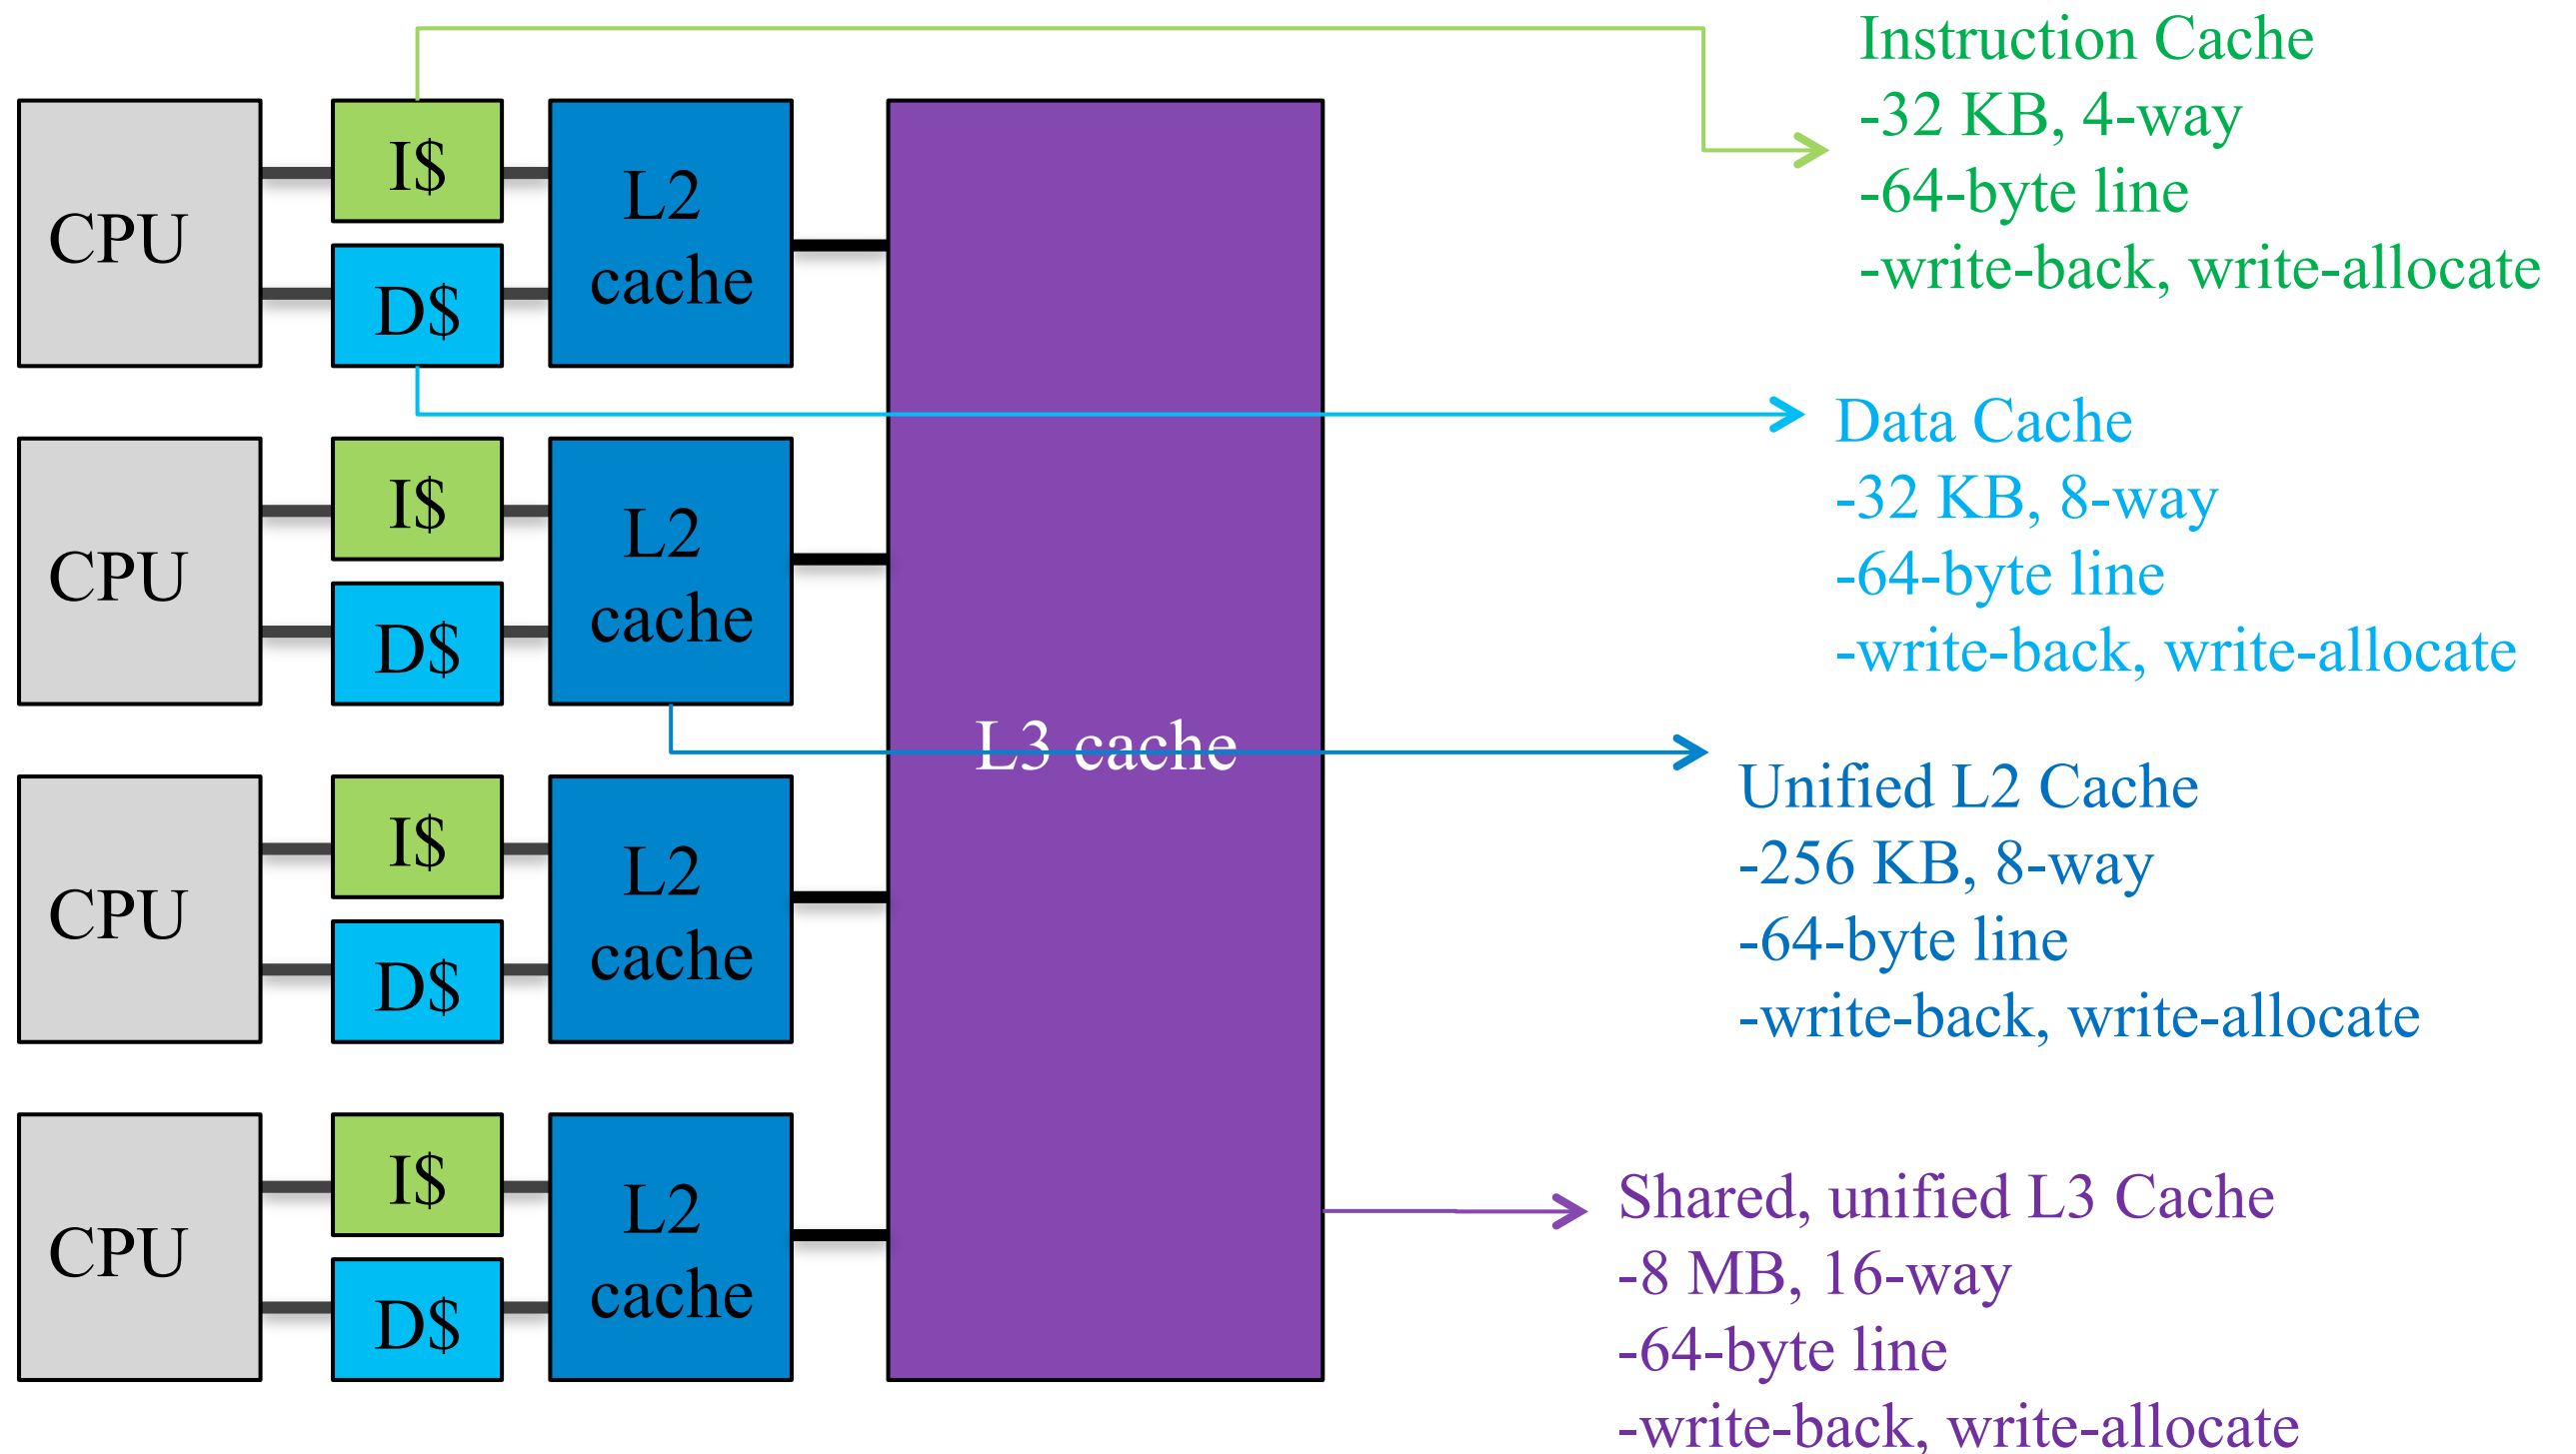
- Keep memory and cache identical?
- **Write-through** => all writes go to both cache and main memory
- **Write-back** => writes go only to cache. Modified cache lines are written back to memory when the line is replaced.
- Make room in cache for store miss?
- **Write-allocate** => on a store miss, bring written line into the cache
- **Write-around** => on a store miss, ignore cache

# The Three C's

---

- Compulsory (or cold-start) misses
  - first access to the data.
- Capacity misses
  - we missed only because the cache isn't big enough.
- Conflict misses
  - we missed because the data maps to the same line as other data that forced it out of the cache.

# Modern Caches





# Key Points

---

# Key Points

---

- Caches give illusion of a **large, cheap** memory with the access time of a fast, expensive memory.

# Key Points

---

- Caches give illusion of a **large, cheap** memory with the access time of a fast, expensive memory.
- Caches take advantage of memory locality, specifically **temporal locality** and **spatial locality**.

# Key Points

---

- Caches give illusion of a **large, cheap** memory with the access time of a fast, expensive memory.
- Caches take advantage of memory locality, specifically **temporal locality** and **spatial locality**.
- Cache design presents many options (block size, cache size, associativity, write policy) that an architect must combine to minimize miss rate and access time to maximize performance